EC527: High Performance Programming with Multicore and GPUs
pthreads Tutorial

NOTE: Depending on your version of pthreads, some of these source files might not compile on your own machine. Try compiling `test_barrier.c` **or** `test_sor.c`**; if you see an error such as** `unknown type name 'pthread_barrier_t'`**, you will need to do these exercises on the Lab machines (accessed as described in the lab_orientation PDF) .**
(Mac OS X users might be able to get these to work if they #include the additional file "apple_pthread_barrier.h")

Here is the template for creating a new thread:

_____

```
int pthread_create(
  pthread_t *thr             // contains new thread's ID
  const pthread_attr_t *attr  // give attributes
  void *(*start_routine)(void) // name of function that the thread will run. It must
                             // have void pointer as its return and parameter values
  void *arg)                 // void pointer for passing parameters
```
_____

Note that there is only one place to pass arguments to the thread, the `void` pointer that is the last parameter of the function. This turns out to be a completely general way to pass parameters, but a perhaps a little non-obvious/unusual. Some of the following is taken from www.faqs.org/docs/learnc/x658.html First let's discuss `void` pointers themselves:

When a variable is declared as being a pointer to type `void` it is known as a generic pointer. Since you cannot have a variable of type `void`, the pointer will not point to any data and therefore cannot be dereferenced. That is, the compiler won't let you get a value from the pointed to location (using "*") because it doesn't know how many bytes to get. A `void` pointer is still a pointer though, but to use it you just have to cast it to another kind of pointer first — hence the term "generic pointer".

A generic pointer is very useful when you want a pointer to point to data of different types at different times, as in when you are creating a generic thread.

_____
test_generic.c

`test_generic.c` uses a `void` pointer (without threads). The cast (`int*`) turns it from "pointer to `void`" into "pointer to `int`", and then "`*(int*)`" tells the compiler you want the value of the `int` pointed to by that. The way to read this is out-to-in. Once you get this basic structure, you can pass anything, as long as it has a defined starting address. (You can even make the address itself the argument to pass by value, as is shown below.)

Task 1. Modify `test_generic.c` so that it used the generic pointer to print out the first three elements of "`float ArrayA`". You should be able to do this in two ways: by adding the offsets (0,1,2) to the pointer and by using array syntax `[0]`, `[1]`, `[2]`. Hint: be careful with your parentheses!

_____
test_create.c

`test_create.c` is the "Hello World!" of thread creation. Read and understand the code. In this example of `pthread_create()` we are not declaring any parameters or special attributes; therefore the values for these fields are `NULL`. Even so, a pointer is passed to the created function so that it knows where to store the thread ID value it generates. It also passes the work function, which denotes where the new thread will begin.

Task 2. Compile and run the code. Record the output. Try it a few times. Does the output change? (It may or may not!)

Task 3. Modify `test_create.c` so that "id" is declared using malloc, and pass the location to store the thread ID information in pointer form (instead of a referenced array).

Task 4. Add a `sleep(3);` line to the `work()` function before the `printf` statement and compile and run. Observe the output when the program is run. What do you see? Why?

Task 5. Remove `sleep(3)` from `work()` and add it to `main()` between the `printf()` and `return()`. Compile, run, and observe the output. What do you see? Why?

What you are seeing is the main thread completing before the child threads can reach the `printf` statement. When the main thread completes, all child threads are automatically killed.

_____
test_join.c (similar to ex3 in the class notes)

`test_join.c` is next. This time, after the threads are created, `main()` walks through the list of thread Ids with the `pthread_join` function, and will successively block on each until that child thread has completed.

Task 6. Add the `sleep(3);` to the child process again, and observe the output. Does it change? Why or why not?

Keep this in mind, because you may need to add joins to the remaining examples to get the expected results.

_____
test_param1.c (similar to ex2b.c in the class notes)

Now let's pass a value to the new thread. Look at `test_param1.c`. We are passing the value of t by casting t to a generic pointer. NOTE: We are not passing the memory location of t, then fetching the value at that location in the new thread. Rather, we are telling the compiler that the value of t is actually a generic pointer, and then the new thread is telling the compiler that no, it is not a memory location, it is actually a value. As long as the type that gets converted to the generic pointer is the same as that to which the generic pointer is converted, this should work.

Task 7. print the value of `threadid` before it is cast back to "`int`" (ignore a compiler warning about `printf` argument type mismatch, if you get one). It should print the same number. Now try passing something really different, like a `signed char` set to some negative number. Does it still work? Why or why not?

_____
test_param2.c (similar to ex2a.c in the class notes)

Look at **test_param2.c**, and observe that you can also pass by reference (in addition to the unorthodox "pass by value" shown in `test_param1.c`). (In `test_param3.c`, we'll pass a pointer to an entire structure to a thread.) Note that the threads cast the generic pointer (now the input) in two ways: as an integer and as a pointer to an integer. When you change the value being pointed to, this change is visible to all other threads.

Task 8. In `test_param2.c`, confirm that threads affect each others' values. Make the code modify `f` before the `printf`, compile and run. Now do the same with `*g`. What happens? Why? Does your output change from run to run?

Task 9. Without changing `NUM_THREADS` or anything in `main()`, try to modify your code so that fewer threads get created.

Task 10. Modify `test_param2.c` to pass an `int` array (the same one) to each thread.
    We want to show that each thread is accessing the same array by having each thread modify a different element, and then printing out the final array. How can you get each thread to use a unique number (without using explicit synchronization)? We'll soon give a way to do this reliably but for now, you might try:
        - performing a hash function, modulo calculation, etc. on the value from pthread_self()
        - declare a global int initially 0, have each thread read its value and then increment it
Can you think of a method that is reliable and will always result in each thread using a unique number? If so, describe it.

_____
test_param3.c (an extension of ex2c.c in the class notes)

Since `pthread_create()` has only one parameter slot, the only way to pass more than one value is by reference. In the previous tasks, all threads were manipulating the same `int` and then the same `int` array. Also, it was hard for them even to tell who they were themselves. (Of course you can look at your own ID, but this doesn't tell you, for example, whether you're "first" or "last".)

Observe that **test_param3.c** passes complex parameters by reference, but also that each one is tailored to its particular thread, with its own ID assigned by `main()`.

Task 11. Compile and run `test_param3.c` to verify this behavior. Now add a 6th thread whose "`sum`" field contains 1000.

_____
test_sync1.c

So far we have "sync'ed" in one way: by joining threads before exiting. This prevented the race condition caused by the master exiting before the child threads. In the last few sections we also amused ourselves by causing various outputs due to having multiple threads modify shared data. In serious programs this is seldom desired. Much more often we want to control when threads execute code which calls or modifies shared data. This can be done using mutexes.

test_sync1.c spawns a thread which cannot reach its `printf` until the main thread releases control of the mutex. This is done by entering a character and pressing return.

Task 12. Comment out the `trylock` check in the child thread and observe the difference in output.

Task 13. Now change `NUM_THREADS` to 2 (and change the `*Messages` initialization), compile and run. Does it work?

---
test_barrier.c

In #13 you observed that sync'ing multiple threads using a single mutex is problematic.

Now look at **`test_barrier.c`**. When a thread reaches a barrier call, it will block until the number of threads specified in the barrier declaration have hit the barrier, at which point all threads will unblock.

Task 14. Compile and run `test_barrier.c`. Observe that all of the "before barrier" messages are printed before the "after barrier" messages. Modify `test_barrier.c` by uncommenting the "`sleep(1)`" statement. Observe the change in behavior. Does the barrier still work? Change the sleep time to be a variable function of taskid. Does the barrier still work?

---
test_sync2.c

Sometimes barriers are too clumsy. That is, we would like threads to execute in a particular sequence. See, e.g., the graph in the lecture notes (file title "L07_ReviewConcurrency.pdf", slide title "Semaphores for Process Synchronization").
`test_sync2.c` implements this graph:

2 & 3	wait for 1 (`main()`)
4	waits for 2
5	waits for 4
6	waits for 3 & 4
7	waits for 5 & 6

The method used is for `main()` (thread 1) to create and lock all of the mutex locks and then create threads 2-7. These threads look at their IDs (passed from `main()`) and go the the corresponding "case." There they wait for the specified lock (2&3 for 1, etc.). When this thread is unlocked by another thread, this thread continues (`break;`). At the end of the thread, it unlocks a lock corresponding to its own ID. This causes other threads to be released. The cascade begins when `main()` unlocks mutex 1.

Note that `test_sync2.c` uses LOCK to wait for an UNLOCK rather than TRY. This blocks and is often more efficient.

Task 15. Compile and run `test_sync2.c` and confirm that the threads run in the order just described. Note that the `join()` loop is executed after most of the threads have exited. Why does this work?

Task 16. Modify `test_sync2.c` to add a thread 8. It waits for threads 6 & 7.

---
test_crit.c

The critical section problem arises when there is a section of code whose simultaneous execution by multiple threads causes an error. **`test_crit.c`** shows the classic example: some threads add to the balance, some subtract. In the end, the balance should stay the same.

Task 17. Compile and run `test_crit.c`. Is the answer correct? That is, is the final balance the same as the beginning? Perhaps you find that it is! (So what's the problem?) Can you change the timing, but not the logic, so that the balance becomes incorrect? How about if you increase the number of threads, to say 10,000? After compiling, run it multiple times to see if you get the same balance each time. Also notice if the printed "`qr_total`" is the same each time.

Task 18. Use a mutex to ensure that it gets the correct final balance each time it is run.

_____

test_sor.c

Successive Over-relaxation (SOR) is a computational technique that can find steady-state solutions for certain systems whose behavior is described by partial differential equations (PDE). A similar computational technique is also used to simulate a PDE system over time. Next week we will look at SOR in detail, in 2 dimensions; this week we have a 1D example.

You are given `test_sor.c`. It declares some variables, and the code for a worker thread is provided.

Task 19. Modify `main()` to create the threads and pass suitable values in their `thread_data` struct, and initialize the `pthread_barrier_t` object.
   Now run it. Is the answer correct? (A block comment shows what the values will be if the system has reached steady-state) How many iterations do you need to get something close to the steady-state solution? Time it to find out how long it takes (using the `time` shell command is okay - but pay attention only to the "real" time, not "sys" or "user").

Task 20. Set `USE_BARRIERS` to 1 and try again. You should be able to get something equally close to steady-state, but with fewer iterations. Time this version to see how long it takes.
   Comparing your findings with `USE_BARRIERS` set to 0 or to 1:
   • Which is faster, and why?
   • Which takes more iterations, and why?
   • Which consistently gives the same answer, and why?